

# AWS Monitoring & Audit: CloudWatch,x-ray

---

## Monitoring

### why monitoring is important ?

- we know how to deploy applications
  - safely
  - automatically
  - using infrastructure as code
  - leveraging the best AWS components
- our application are deployed, and our users don't care how we did it
- our users only care that the application is working
  - application latency: will it increase over time ?
  - application outages : customer experience should not be degraded
  - users contacting the IT department or complaining [抱怨] is not a good outcome
  - troubleshooting and remediation [补救措施]
- Internal monitoring
  - can we prevent issues before they happen ?
  - performance and cost
  - trends (scaling patterns)
  - learning and improvement

## Monitoring in AWS

- AWS CloudWatch:
  - Metrics : collect and track key metrics
  - Logs : collect, monitor, analyze and store log files
  - Events : send notifications when certain events happens in your AWS
  - Alarms : react in real-time to metrics /events
- AWS X-Ray
  - Troubleshooting application performance and errors
  - distributed tracing of microservices
- AWS CloudTrail
  - internal [内部的] monitoring of API calls being made
  - audit [审核] changes to AWS Resources by your users

## CloudWatch

### Metrics

- CloudWatch provides metrics for every services in AWS
- Metrics is a variable to monitor (CPUUtilization ,NetworkIn..)
- Metrics belong to namespaces
- Dimension [维度] is an attribute of a metric (instance id,environment, etc,,,)

- up to 30 dimensions per metric
- Metrics have timestamps [时间戳]
- can create CloudWatch dashboards of metrics

## EC2 Detailed monitoring

- EC2 instance metrics have metrics "every 5 minutes"
- with detailed monitoring (for a cost), you get data "every 1 minute"
- use detailed monitoring if you want to scale faster for your ASG
- the AWS free tier allows us to have 10 detailed monitoring metrics
- Note: EC2 Memory usage is by default not pushed (must be pushed from inside the instance as a custom metric)

## Custom Metrics

- possibility to define and send your own custom metrics to CloudWatch
- example: memory(RAM) usage, disk space, number of logged in users..
- use API call **PutMetricData**
- ability to use dimensions [维度](attributes) to segment [细分] metrics
  - instance.id
  - environment.name
- metric resolution (**StorageResolution** API parameter – two possible values):
  - standard : 1 minute (60 seconds)
  - high resolution: 1/5/10/30 seconds – higher cost
- important : accepts metric data points 2 weeks in the past and 2 hours in the future (make sure to configure your EC2 instance time correctly)

## Logs

- log groups : arbitrary [随机] name, usually representing [代表] an application
- log streams : instances within application / log files / containers
- can define log expiration policies (never expire, 1 day to 10 years)
- CloudWatch logs can send logs to :
  - amazon S3 (exports)
  - Kinesis Data Streams
  - Kinesis Data Firehose
  - AWS Lambda
  - OpenSearch
- logs are encrypted by default
- can setup KMS-based encryption with your own keys

## Source

- SDK, CloudWatch Logs Agent, CloudWatch Unified [统一] Agent
- Elastic Beanstalk : collection of logs from application
- ECS : collection from containers
- AWS Lambda: collection from function logs
- VPC Flow Logs : VPC specific logs

- API Gateway
- CloudTrail based on filter
- Route53 :Log DNS queries

## logs insights

The screenshot shows the CloudWatch Logs Insights interface. At the top, there's a search bar with 'application.log' selected and a 'Clear' button. Below it is a query editor with the following code:

```
1 fields @timestamp, @message
2 | sort @timestamp desc
3 | limit 20
```

Next to the editor are buttons for 'Run query', 'Save', and 'History'. A note says 'Queries are allowed to run for up to 15 minutes.' To the right of the editor are buttons for 'Change the time range here.', 'Discovered Fields in your log groups.', and 'Fields' (with a dropdown menu). Below the editor is a time range selector showing '2021-11-09 (06:40:02) > 2021-11-09 (06:55:17)'.

The main area displays a histogram titled 'Showing 20 of 10,197 records matched' with a note '10,197 records (2.3 MB) scanned in 3.3s at 3,091 records/s (714.9 kB/s)'. Below the histogram is a table with two rows of log entries:

#	@timestamp	@message
▶ 1	2021-11-09T06:54:17.62...	{"Severity": "INFO", "message": "This is where the message detail would go", "IP Address": "10.30.86.98", "Timestamp": "2021-11-09T11:06:54.176Z"}
▶ 2	2021-11-09T06:54:13.38...	{"Severity": "INFO", "message": "This is where the message detail would go", "IP Address": "192.168.0.43", "Timestamp": "2021-11-09T11:06:54.133Z"}

Below the table is a link: <https://mng.workshop.aws/operations-2022/detect/cwlogs.html>.

This screenshot shows the 'Sample queries' section of the CloudWatch Logs Insights interface. It includes a 'Learn more' link and a sidebar with 'Common queries' and '25 most recently added log events'.

**25 most recently added log events:**

```
fields @timestamp, @message
| sort @timestamp desc
| limit 25
```

**Number of exceptions logged every 5 minutes:**

```
filter @message like /Exception/
| stats count(*) as exceptionCount by
bin(5m)
| sort exceptionCount desc
```

**List of log events that are not exceptions:**

```
fields @message
| filter @message not like /Exception/
```

- search and analyze log data stored in CloudWatch Logs
- example: find a specific IP inside a log, count occurrences [发生] of "ERROR" in your logs
- provides a purpose-built [构建] query language
  - automatically discovers fields from AWS services and JSON log events
  - fetch desired [想要的] event fields, filter based on conditions, calculate aggregate [总计] statistics [统计], sort events, limit number of events..
  - can save queries and add them to CloudWatch Dashboards
- can query multiple Log Groups in different AWS accounts

- it's a query engine, not a real-time engine

## S3 Export



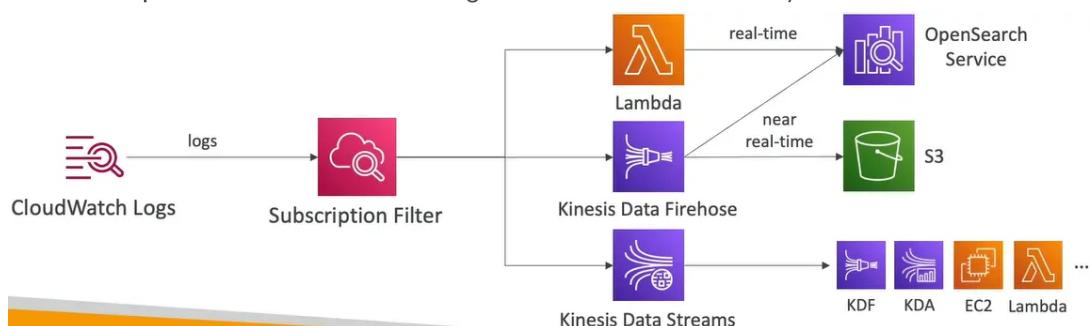
## CloudWatch Logs

Amazon S3

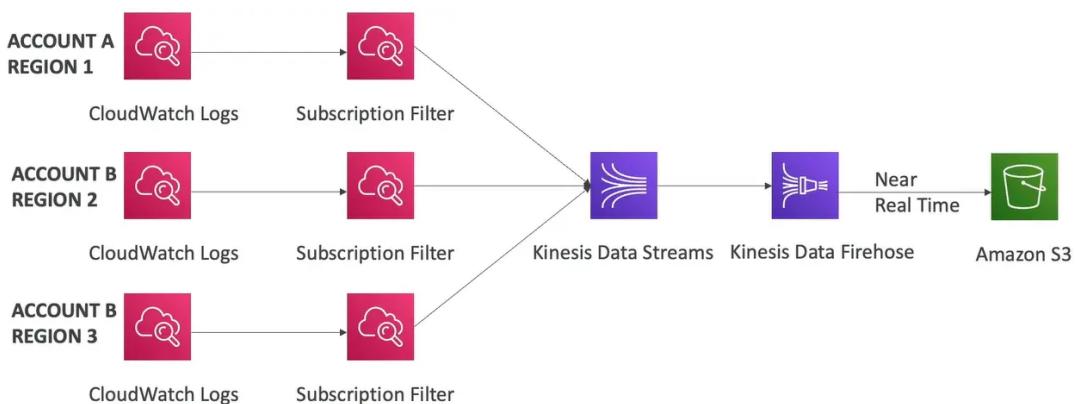
- log data can take up to 12 hours to become available for export
  - the API call is `CreateExportTask`
  - not near-real time or real-time... use logs Subscriptions instead

## Logs Subscriptions

- get real-time log events from CloudWatch logs for processing and analysis
  - send to Kinesis Data Streams, Kinesis Data Firehose, or Lambda
  - **Subscription Filter** – filter which logs are events delivered to your destination

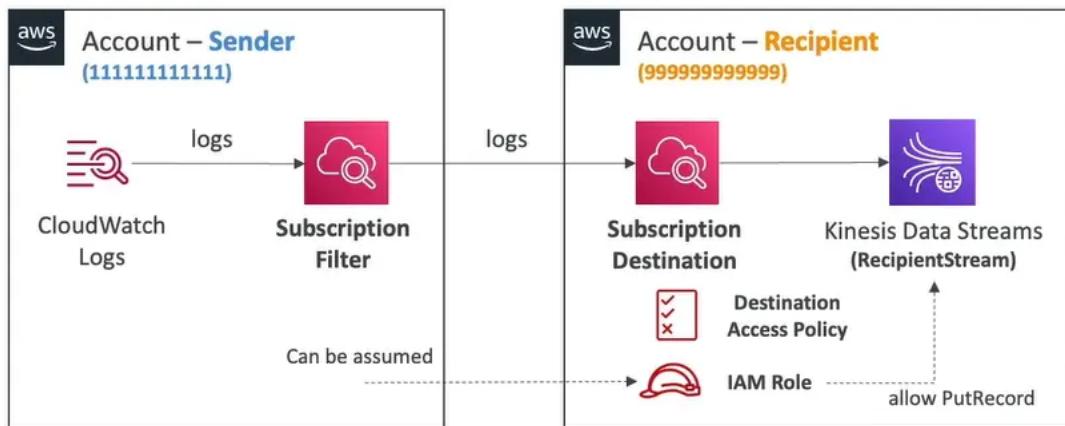


logs aggregation [聚合] multi-account & multi region



## CloudWatch Logs Subscriptions

- Cross-Account Subscription – send log event to resource in a different AWS account(KDS,KDF)



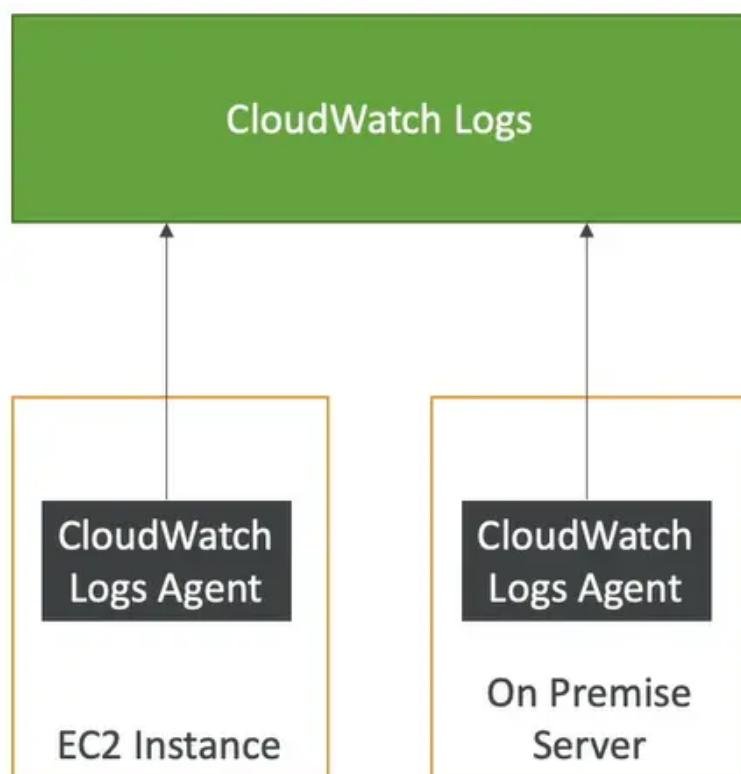
```

{
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "kinesis:PutRecord",
      "Resource": "arn:aws:kinesis:us-east-1:
999999999999:stream/RecipientStream"
    }
  ]
}
  
```

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "AWS": "111111111111"
      },
      "Action": "logs:PutSubscriptionFilter",
      "Resource": "arn:aws:logs:us-east-1:999999999999:
destination:testDestination"
    }
  ]
}
  
```

### Logs for EC2



- by default , no logs from your EC2 machine will go to CloudWatch
- you need to run a CloudWatch agent on EC2 to push the log files you want
- make sure IAM permissions are correct
- the CloudWatch log agent can be setup on-premises too

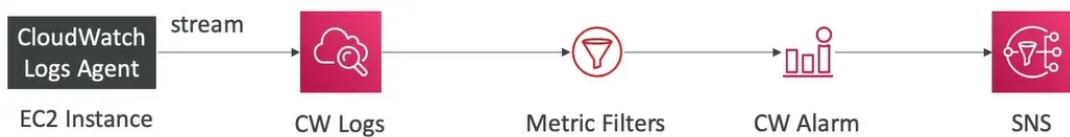
## CloudWatch Logs Agent & Unified Agent

- for virtual servers (EC2 instance, on-premise servers..)
- CloudWatch Logs Agent
  - old version of the agent
  - can only send to CloudWatch Logs
- CloudWatch Unified Agent
  - Collect additional system-level metrics such as RAM, process, etc.,,
  - collect logs to send CloudWatch Logs
  - centralized [统一] configuration using SSM Parameter Store

## CloudWatch Unified Agent –Metrics

- collected directly on your Linux server / EC2 instance
- CPU (active, guest, idle [闲置], system, user, steal)
- disk metrics (free, used, total, ), disk IO (writes, reads, bytes, iops)
- RAM (free, inactive, used, total, cached)
- Netstat (number of TCP and UDP connections, net packets, bytes)
- Processes (total, dead, blocked, idle, running, sleep)
- Swap Space [交换空间] (free, used, used %)
- Remember : out-of-the box metrics for EC2 –disk, CPU, network (high level)

## CloudWatch Logs Metric Filter



- CloudWatch logs can use filter expressions
  - for example, find a specific IP inside of a log
  - or count occurrences of "ERROR" in your logs
  - metric filters can be used to trigger [触发] alarms
- filter do not retroactively [追溯] filter data, filters only publish the metric data points for events that happened after the filter was created
- ability to specify up to 3 Dimensions [维度] for the Metric Filter (optional)

## Alarm

- alarms are used to trigger notification for any metric
- various options (sampling, %, max, min, etc..)
- alarm states:
  - ok
  - insufficient [不足] \_data
  - alarm
- period:

- length of time in seconds to evaluate the metric
- high resolution custom metrics: 10sec, 30sec or multiples of 60sec

## targets

- stop, terminate, reboot, or recover an EC2 instance
- trigger auto scaling action
- send notification to SNS (from which you can do pretty much anything)



Amazon EC2



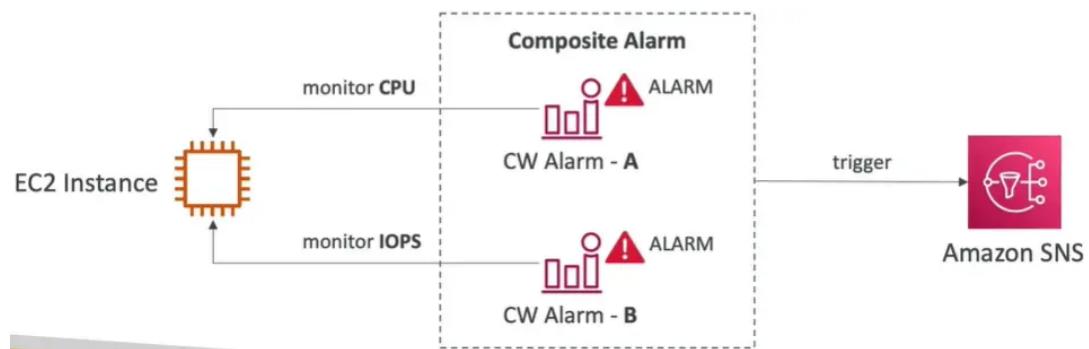
EC2 Auto Scaling



Amazon SNS

## Composite [合成的] Alarms

- cloudwatch alarms are on a single metric
- composite alarms are monitoring the states of multiple other alarms
- AND and OR conditions
- helpful to reduce "alarm noise" by creating complex composite alarms



## EC2 instance recovery

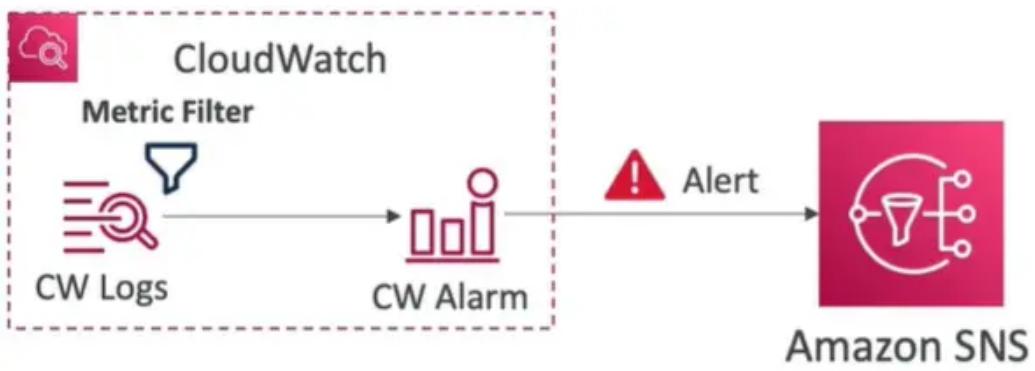
- status check:
  - instance status = check the EC2 VM
  - system status = check the underlying hardware



- Recovery : same private,public,elastic,IP,metadata,placement group

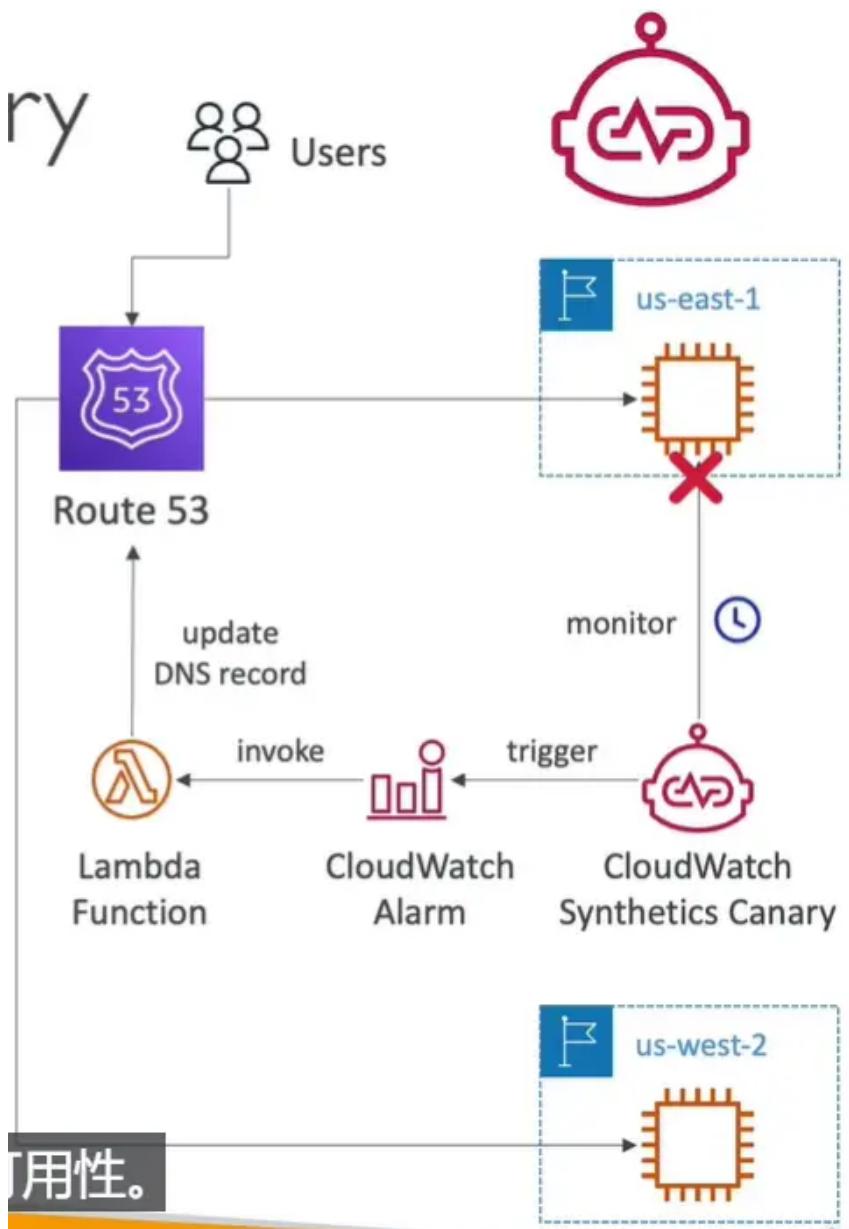
## good to know

- alarms can be created based on CloudWatch logs metrics filters



- to test alarms and notifications, set the alarm state to Alarm using CLI  
`aws cloudwatch set-alarm-state --alarm-name "myalarm" --state-value ALARM --state-reason "testing purposes"`

## Synthetics [合成的] Canary



- configurable script that monitor your APIs, URLs, Websites...
- reproduce what your customers do programmatically to find issues before customers are impacted [受影响的]

- checks the availability and latency of your endpoints and can store load time data and screenshots of the UI
- integration with CloudWatch Alarms
- scripts written in Node.js or Python
- programmatic access to a headless Google Chrome browser
- can run once or on a regular schedule

## Synthetics Canary Blueprints

- heartbeat monitor – load URL, store screenshot and an HTTP archive file
- API Canary – test basic read and write functions of REST APIs
- broken link checker – check all links inside the URL that you are testing
- visual monitoring – compare a screenshot taken during a canary run with a baseline screenshot
- canary recorder – used with CloudWatch Synthetics Recorder (record your actions on a website and automatically generates a script for that)
- GUI Workflow Builder – verifies that actions can be taken on your webpage (e.g. test a webpage with a login form)

## Amazon EventBridge (formerly CloudWatch Events)

- Schedule : Cron jobs (schedule scripts)

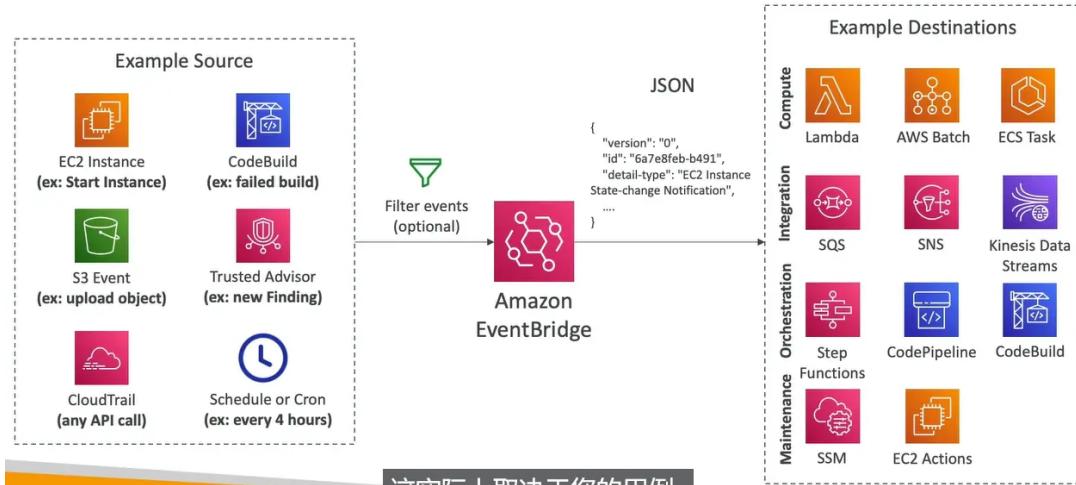


- Event pattern : event rules to react to a service doing something



- trigger Lambda functions, send SQS / SNS message...

## rules



## amazon eventbridge



- event buses can be accessed by other AWS accounts using Resource-based Policies
- you can archive events (all / filter) sent to an event bus (indefinitely or set period)
- ability to replay archived events

## Schema Registry [模式注册表]

The screenshot shows the AWS Schema Registry interface for a schema named `aws.codepipeline@CodePipelineActionExecutionStateChange`. The schema details are as follows:

Schema name	Last modified	Schema ARN
<code>aws.codepipeline@CodePipelineActionExecutionStateChange</code>	Dec 1, 2019, 12:11 AM GMT	-
Description		Schema registry aws.events Number of versions 1 Schema type OpenAPI 3.0
Schema for event type	<code>CodePipelineActionExecutionStateChange</code> , published by AWS service <code>aws.codepipeline</code>	

**Version 1** (Created on Dec 1, 2019, 12:11 AM GMT)

Action ▾ Download code bindings

```

1 {
2   "openapi": "3.0.0",
3   "info": {
4     "version": "1.0.0",
5     "title": "CodePipelineActionExecutionStateChange"
6   },
7   "paths": {},
8   "components": {
9     "schemas": {
10       "AWSEvent": {
11         "type": "object",
12         "properties": {
13           "state": {
14             "type": "string",
15             "enum": [
16               "SUCCEEDED",
17               "FAILED",
18               "TIMED_OUT",
19               "SKIPPED"
20             ]
21           }
22         }
23       }
24     }
25   }
26 }
```

- EventBridge can analyze the events in your bus and infer [推断] the schema
- the schema registry allows you to generate code for your application, that will know in advance [提前] how data is structured in the event bus
- schema can be versioned

## Resource-based Policy

The diagram shows a Resource-based Policy allowing events from another AWS account to an EventBridge Bus.

**Policy Document:**

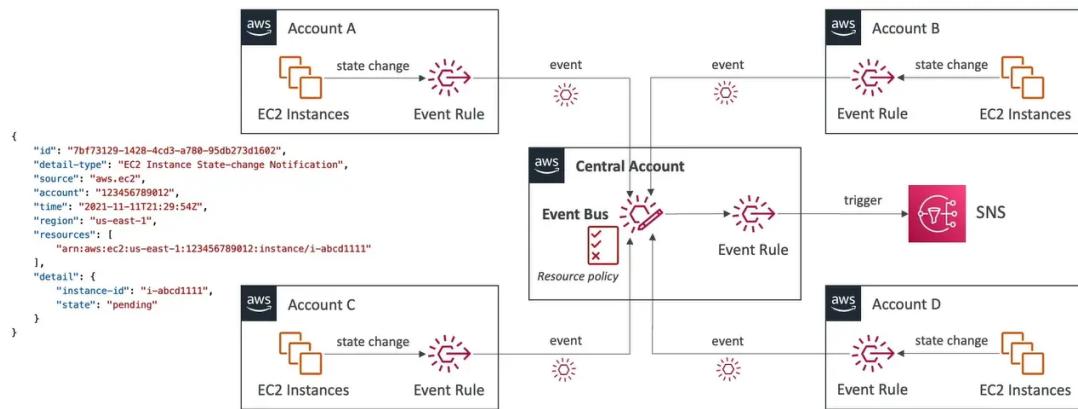
```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "events:PutEvents",
      "Principal": { "AWS": "111122223333" },
      "Resource": "arn:aws:events:us-east-1:123456789012:central-event-bus"
    }
  ]
}
```

**Allow events from another AWS account**

**EventBridge Bus (central-event-bus)** (AWS Account 123456789012) has a PutEvents permission to the **Lambda function** (AWS Account 111122223333).

- manage permissions for a specific Event Bus
- example: allow / deny events from another AWS account or AWS region
- use case : aggregate [聚合] all events from your AWS Organization in a single AWS account or AWS region

## multi-account aggregation



## AWS X-Ray

- debugging in production, the good old way:
  - test locally
  - add log statement everywhere
  - re-deploy in production
- log formats differ across applications using CloudWatch and analytics is hard
- debugging: monolith [单体] "easy", distributed services "hard"
- no common views of your entire architecture

## visual analysis of our applications



## advantages

- troubleshooting performance (bottlenecks [瓶颈])
- understanding dependencies in a microservice architecture
- pinpoint [查明] service issues
- review request behavior
- find errors and exceptions
- are we meeting time SLA ?
- where am I throttled [限制]?
- identify users that are impacted

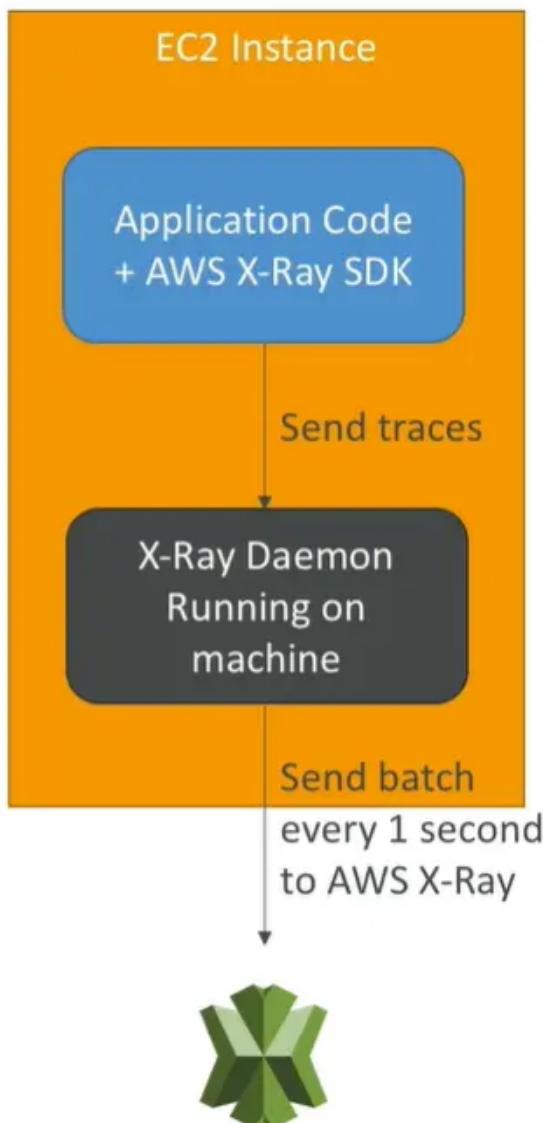
## compatibility

- aws lambda
- elastic beanstalk
- ecs
- elb
- api gateway
- ec2 instance or any application server (even on premise [前提])

## Leverages Tracing

- tracing is an end to end way to following a "request"
- each component dealing with the request adds its own "trace"
- tracing is made of segments [段] (+sub segments)
- annotations [注释] can be added to traces to provide extra-information
- ability to trace:
  - every request
  - sample request (as a % for example or a rate per minute)
- X-Ray Security
  - IAM for authorization
  - KMS for encryption at rest

how to enable it ?



1. your code (java,python,go,node.js,.net) must import the AWS X-Ray SDK

- a. very little code modification needed
- b. the application SDK will then capture
  - i. calls to AWS services
  - ii. HTTP / HTTPS requests
  - iii. database calls (mySQL,PostgreSQL,DynamoDB)
  - iv. queue calls (SQS)
- 2. install the X-Ray daemon [守护进程] or enable X-Ray AWS Integration
  - a. X-Ray daemon works as a low level UDP packet interceptor [拦截器] (Linux /windows/mac)
  - b. aws lambda /other AWS services already run the X-Ray daemon for you
  - c. each application must have the IAM rights to write data to X-Ray

## the X-Ray magic



- X-Ray service collects data from all the different services
- service map is computed from all the segments and traces
- X-Ray is graphical [图形化的], so even no technical people can help troubleshoot

## Troubleshooting

- if X-Ray is not working on EC2
  - ensure the EC2 IAM Role has the proper [正确的] permissions
  - ensure the EC2 instance is running the X-Ray Daemon
- to enable on AWS Lambda
  - ensure it has an IAM execution role with proper policy (AWSX-RayWriteOnlyAccess)
  - ensure that X-Ray is imported in the code
  - enable Lambda X-Ray Active Tracing

## X-Ray Instrumentation [检测] in your code

## Example for Node.js & Express

```
var app = express();

var AWSXRay = require('aws-xray-sdk');
app.use(AWSXRay.express.openSegment('MyApp'));

app.get('/', function (req, res) {
  res.render('index');
});

app.use(AWSXRay.express.closeSegment());
```

- instrumentation means the measure of product's performance, diagnose [诊断] errors, and to write trace information
- to instrument your application code, you use the X-Ray SDK
- many SDK require only configuration changes
- you can modify your application code to customize and annotation the data that the SDK sends to X-Ray, using interceptors, filters, handlers, middleware...

## Concepts [概念]

- segments: each application / service will send them
- subsegments: if you need more details in your segment
- trace: segments collected together to form an end-to-end trace
- sampling[采样]: decrease the amount of requests sent to X-Ray, reduce cost
- annotations: key value pairs used to index traces and use with filters
- metadata: key value pairs, not indexed, not used for searching
- the X-Ray daemon / agent has a config to send traces cross account
  - make sure the IAM permissions are correct – the agent will assume the role
  - this allows to have a central account for all your application tracing

## Sampling Rules

- with sampling rules, you control the amount of data that you record
- you can modify sampling rules without changing your code
- by default ,the X-Ray SDK records the first request **each second**, and 5 percent of any additional requests
- **one request per second** is the reservoir, which ensures that at least one trace is recorded each second as long the service is serving requests
- **5 percent** is the rate at which additional requests beyond the reservoir size are sampled

## Custom Sampling Rules

#### Example Higher minimum rate for POSTs

- Rule name – POST minimum
- Priority – 100
- Reservoir – 10
- Rate – 0.10
- Service name – \*
- Service type – \*
- Host – \*
- HTTP method – POST
- URL path – \*
- Resource ARN – \*

#### Example Debugging rule to trace all requests for a problematic route

- A high-priority rule applied temporarily for debugging.
- Rule name – DEBUG – history updates
  - Priority – 1
  - Reservoir – 1
  - Rate – 1
  - Service name – Scorekeep
  - Service type – \*
  - Host – \*
  - HTTP method – PUT
  - URL path – /history/\*
  - Resource ARN – \*

- you can create your own rules with the reservoir and rate

## API

### X-Ray Write APIs (used by the X-Ray daemon)

```
"Effect": "Allow",
"Action": [
    "xray:PutTraceSegments",
    "xray:PutTelemetryRecords",
    "xray:GetSamplingRules",
    "xray:GetSamplingTargets",
    "xray:GetSamplingStatisticSummaries"
],
"Resource": [
    "*"
]
```

[arn:aws:iam::aws:policy/AWSXrayWriteOnlyAccess](#)

- PutTraceSegments: uploads segment documents to AWS X-Ray
- PutTelemetryRecords : used by the AWS X-Ray daemon to upload telemetry [遥感数据].
  - SegmentsReceivedCount.
  - SegmentsRejectedCount,BackendConnectionErrors..
- GetSamplingRules: retrieve [检索] all sampling rules (to know what/when to send)
- GetSamplingTargets & GetSamplingStatisticSummaries: advanced
- the X-Ray daemon needs to have an IAM policy authorizing the correct API calls to function correctly

### X-Ray Read APIs – continued

```

"Effect": "Allow",
"Action": [
    "xray:GetSamplingRules",
    "xray:GetSamplingTargets",
    "xray:GetSamplingStatisticSummaries",
    "xray:BatchGetTraces",
    "xray:GetServiceGraph",
    "xray:GetTraceGraph",
    "xray:GetTraceSummaries",
    "xray:GetGroups",
    "xray:GetGroup",
    "xray:GetTimeSeriesServiceStatistics"
],
"Resource": [
    "*"
]

```

- GetServiceGraph: main graph
- BatchGetTraces: retrieves a list of traces specified by ID. Each trace is a collection of segment documents that originates from a single request
- GetTraceSummaries: retrieves IDs and annotations for traces available for a specified time frame using an optional filter. to get the full traces, pass the trace IDs to BatchGetTraces.
- GetTraceGraph : retrieves a service graph for one or more specific trace IDs

## X-Ray with Beanstalk

- AWS Elastic Beanstalk platforms include the X-Ray daemon
- you can run the daemon by setting an option in the Elastic Beanstalk console or with a configuration file (in. ebextensions /xray-daemon.config)

```

option_settings:
  aws:elasticbeanstalk:xray:
    XRayEnabled: true

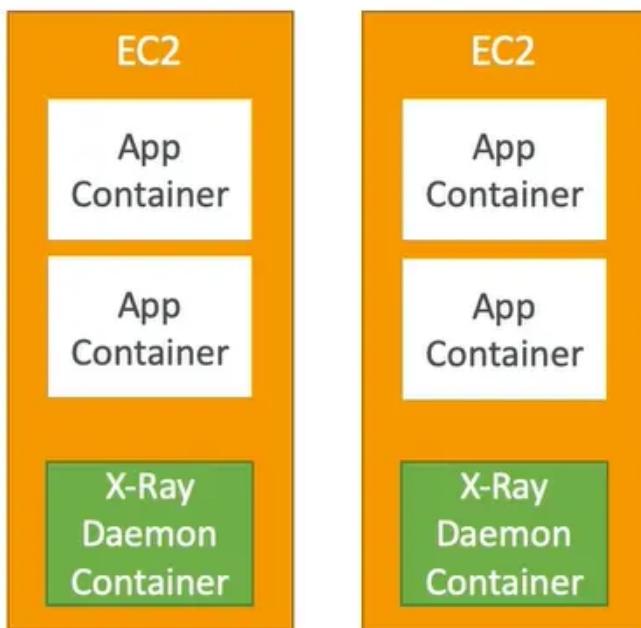
```

- make sure to give your instance profile the correct IAM permissions so that the X-Ray daemon can function correctly
- then make sure your application code is instrumented with the X-Ray SDK
- Note : the X-Ray daemon is not provided for Multicontainer Docker

## ECS + X-Ray integration options

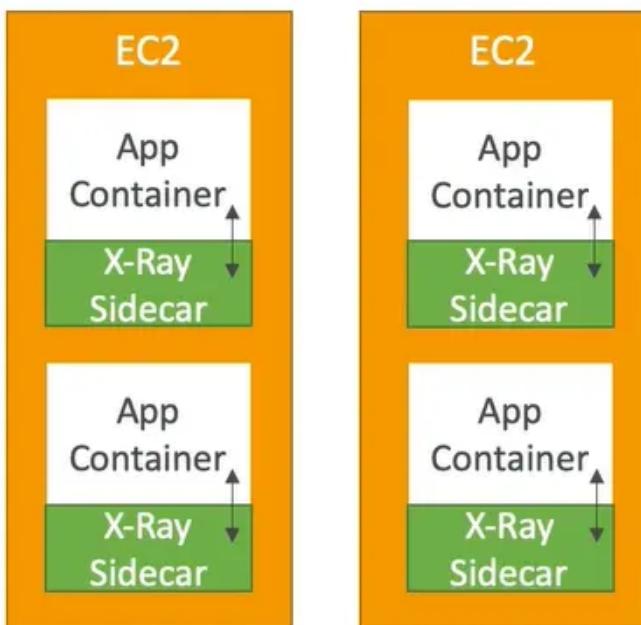
## ECS Cluster

### X-Ray Container as a Daemon



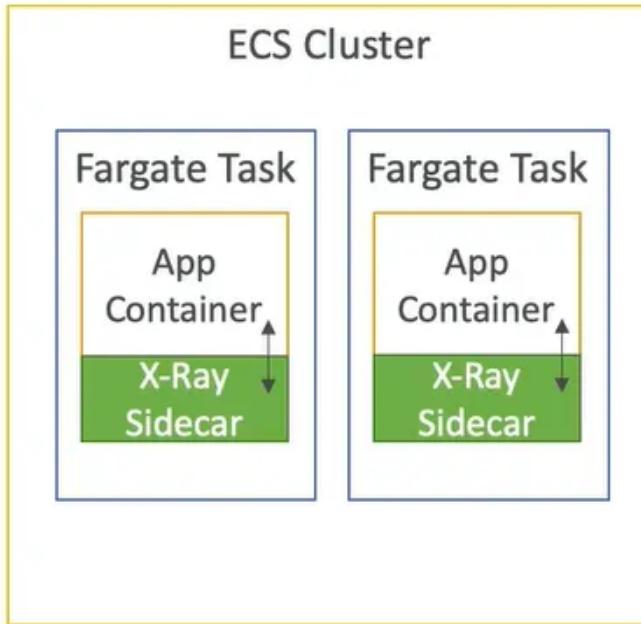
## ECS Cluster

### X-Ray Container as a “Side Car”



# Fargate Cluster

## X-Ray Container as a “Side Car”



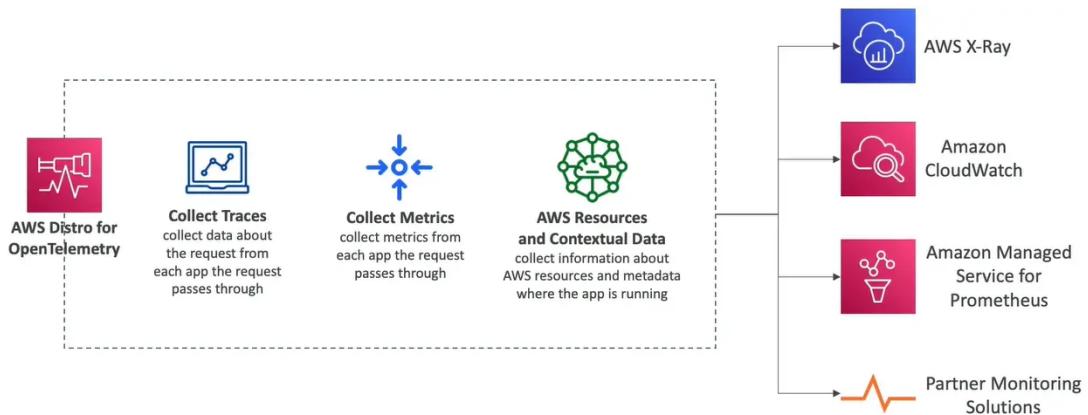
## ECS +X-Ray : Example Task Definition

```
{
  "name": "xray-daemon",
  "image": "123456789012.dkr.ecr.us-east-2.amazonaws.com/xray-daemon",
  "cpu": 32,
  "memoryReservation": 256,
  "portMappings": [
    {
      "hostPort": 0,
      "containerPort": 2000,
      "protocol": "udp"
    }
  ],
  {
    "name": "scorekeep-api",
    "image": "123456789012.dkr.ecr.us-east-2.amazonaws.com/scorekeep-api",
    "cpu": 192,
    "memoryReservation": 512,
    "environment": [
      { "name": "AWS_REGION", "value": "us-east-2" },
      { "name": "NOTIFICATION_TOPIC", "value": "arn:aws:sns:us-east-2:123456789012:scorekeep-notifications" },
      { "name": "AWS_XRAY_DAEMON_ADDRESS", "value": "xray-daemon:2000" }
    ],
    "portMappings": [
      {
        "hostPort": 5000,
        "containerPort": 5000
      }
    ],
    "links": [
      "xray-daemon"
    ]
  }
}
```

## AWS Distro for Open Telemetry

- secure, production-ready AWS-support distribution of the open-source project Open Telemetry project
- provides a single set of APIs, libraries, agents, and collector services
- collects distributed traces and metrics from your apps
- collects metadata from your AWS resources and services
- Auto-instrumentation Agents to collect traces without changing your code
- send traces and metrics to multiple AWS services and partner solutions
  - X-Ray, CloudWatch, Prometheus..

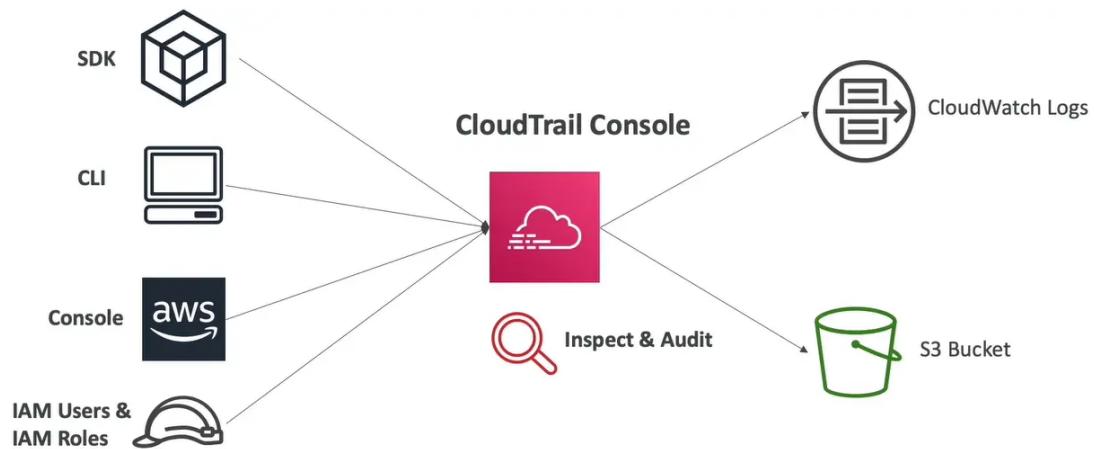
- instrument your apps running on AWS (eg. EC2, ECS, EKS, Fargate, Lambda) as well as on-premises
- migrate from X-Ray to AWS Distro for Telemetry if you want to standardize with open-source APIs from Telemetry or send traces to multiple destinations simultaneously



## AWS CloudTrail

- provides governance [治理], compliance [遵守] and audit [审计] for your AWS Account
- CloudTrail is enabled by default
- get an history of events / API calls made within your AWS Account by
  - console
  - sdk
  - cli
  - aws service
- can put logs from CloudTrail into CloudWatch Logs or S3
- A trail can be applied to All Region (default) or a single Region
- if a resource is deleted in AWS, investigate [调查] CloudTrail first !

## CloudTrail Diagram



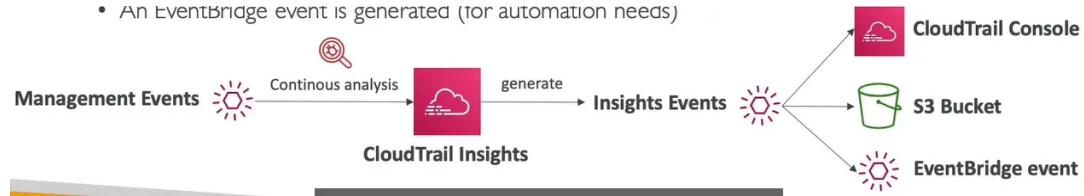
## CloudTrail Events

- Management Events
  - operations that are performed on resource in your AWS account
  - examples:
    - configuring security (IAM AttachRolePolicy)
    - configuring rules for routing data (Amazon EC2 CreateSubnet)
    - setting up logging (AWS CloudTrail CreateTrail)
  - by default, trails are configured to log management events

- can separate Read Events (that don't modify resource) from Write Events (that may modify resources)
- Data Event
  - by default, data events are not logged (because high volume operations)
  - Amazon S3 object-level activity (ex: GetObject,DeleteObject,PutObject): can separate Read and Write Events
  - AWS Lambda function execution activity (the Invoke [调用] API)
- CloudTrail Insights Events

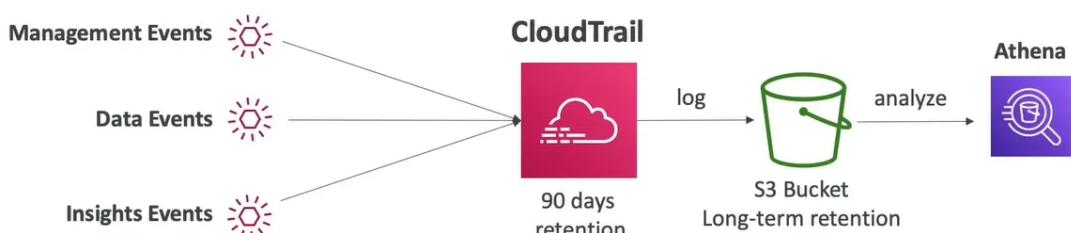
## CloudTrail Insights

- enable CloudTrail Insights to detect [探测] unusual activity in your account
  - inaccurate [不准确] resource provisioning
  - hitting service limits
  - bursts of AWS IAM actions
  - gaps in periodic maintenance activity
- CloudTrail insights analyze normal management events to create a baseline
- and then continuously analyze write events to detect unusual patterns
  - anomalies appear in the CloudTrail console
  - event is sent to Amazon S3
  - an EventBridge event is generated (for automation needs)
    - An EventBridge event is generated (for automation needs)

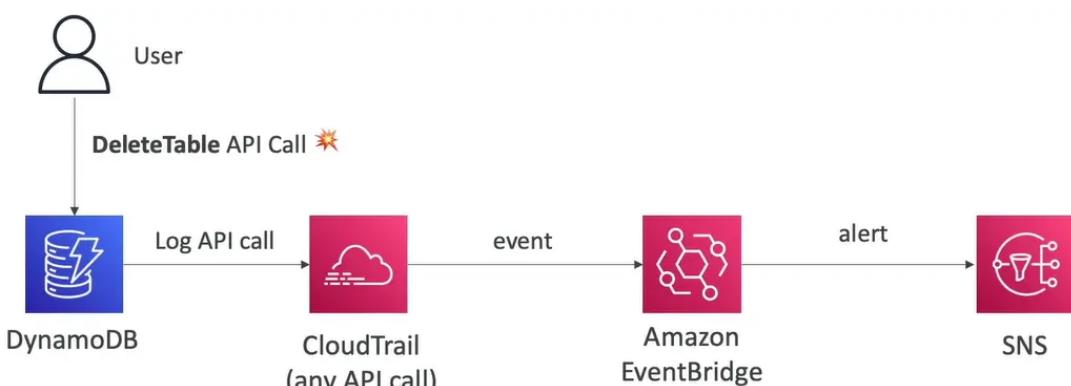


## Event Retention

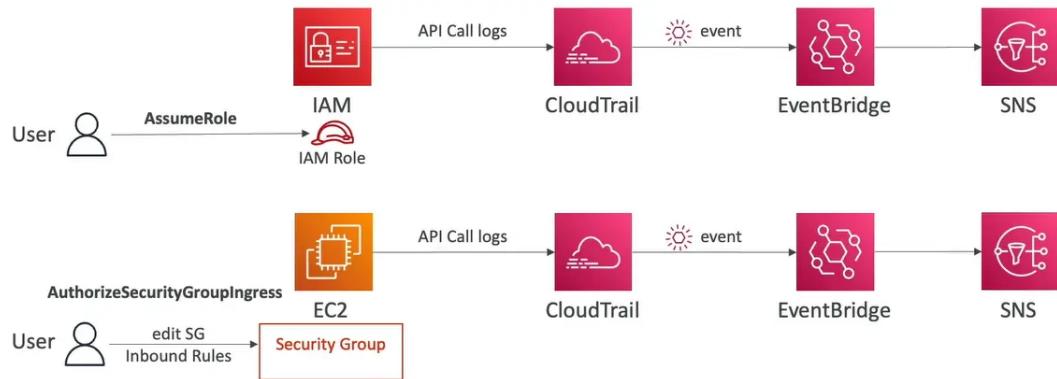
- events are stored for 90 days in CloudTrail
- to keep events beyond this period, log them to S3 and use Athena



## Amazon EventBridge – Intercept API Calls



## Amazon EventBridge – CloudTrail



## CloudTrail vs CloudWatch vs X-Ray

- CloudTrail
  - audit API calls made by users/services /AWS console
  - useful to detect unauthorized calls or root cause of changes
- CloudWatch
  - CloudWatch Metrics over time for monitoring
  - CloudWatch logs for storing application log
  - CloudWatch Alarms to send notification in case of unexpected metrics
- X-Ray:
  - automated Trace Analysis & Central Service Map Visualization
  - latency, Errors and Fault analysis
  - request stracking across distributed system