

AWS Other Services

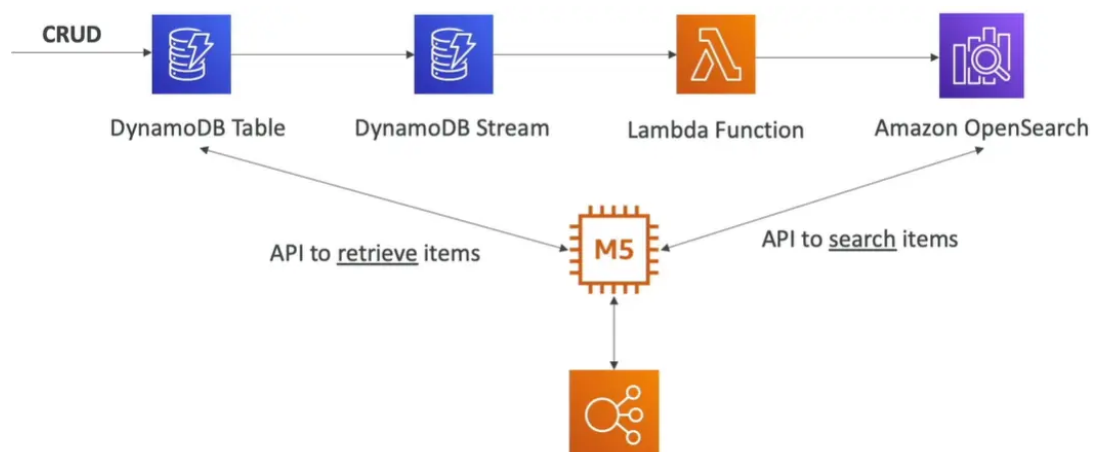
AWS SES – Simple Email Service

- send emails to people using
 - SMTP interface
 - or AWS SDK
- ability to receive email. Integrate with
 - S3
 - SNS
 - Lambda
- Integrated with IAM for allowing to send emails



Amazon OpenSearch Service

- Amazon OpenSearch is successor to Amazon ElasticSearch
- In DynamoDB, queries only exist by primary key or indexes...
- With OpenSearch, you can search an field, even partially matches
- it's common to use OpenSearch as a complement [补充] to another database
- two modes : managed cluster or serverless cluster
- does not natively support SQL (can be enabled via a plugin)
- ingestion [摄入] from Kinesis Data Firehose, AWS IoT, and CloudWatch Logs
- security through Cognito & IAM, KMS encryption, TLS
- comes with OpenSearch Dashboards (visualization)

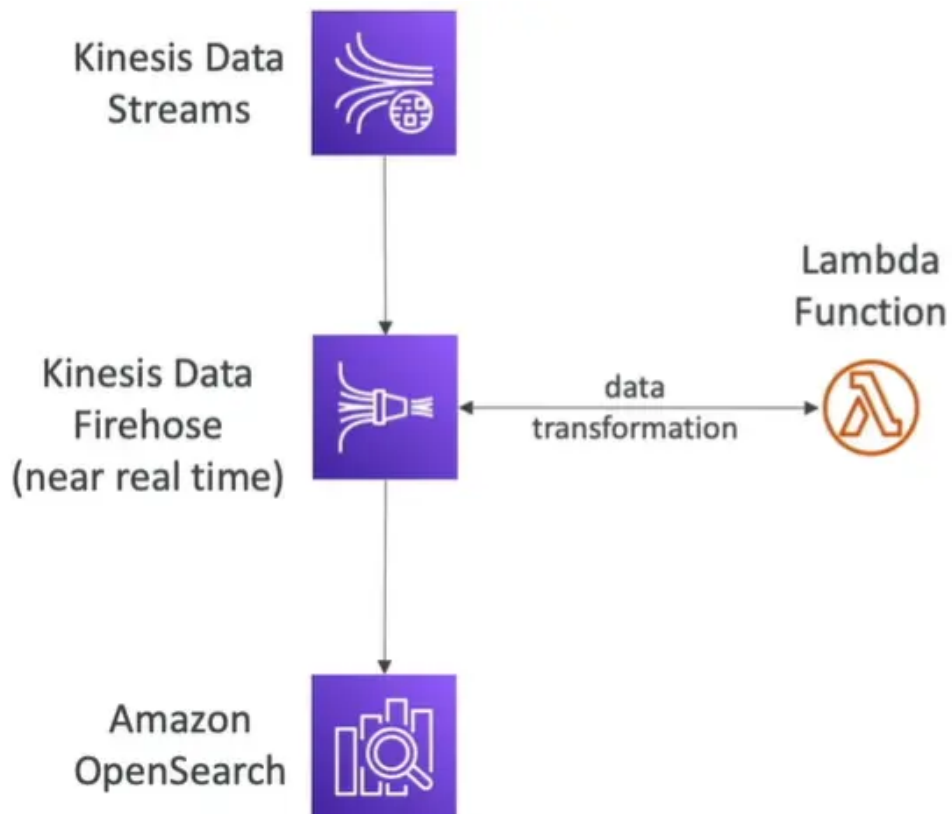
Open Search patterns DynamoDB

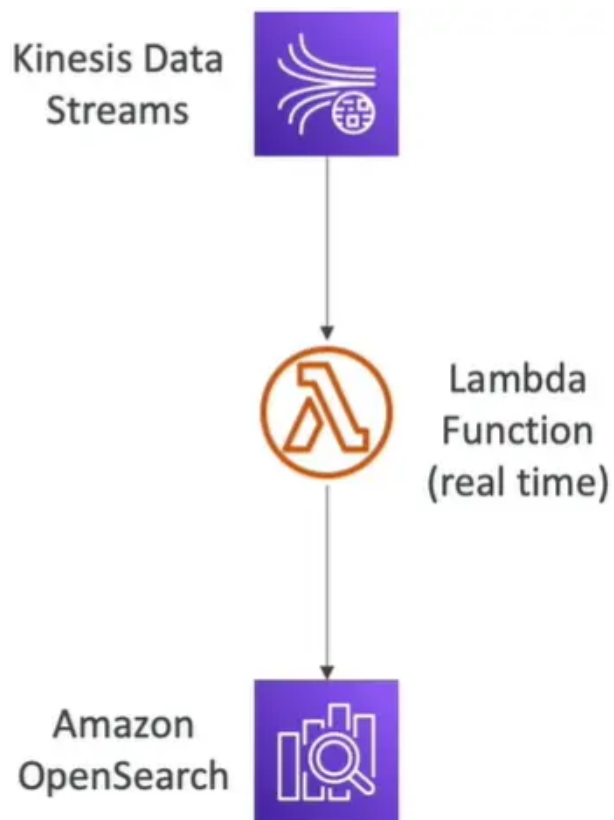


OpenSearch patterns CloudWatch Logs



OpenSearch patterns Kinesis Data Stream & Kinesis Data Firehose





Amazon Athena



- serverless query service to analyze data stored in Amazon S3
- uses standard SQL language to query the files (built on Presto)
- supports CSV, JSON, ORC, Avro and Parquet
- Pricing: \$ 5.00 per TB of data scanned

- commonly used with Amazon QuickSight for reporting / dashboards
- **use cases** : business intelligence / analytics / reporting , analyze & query VPC Flow Logs, ELB logs, CloudTrail trails,etc ...
- Exam Tip : analyze data in S3 using serverless SQL, use Athena

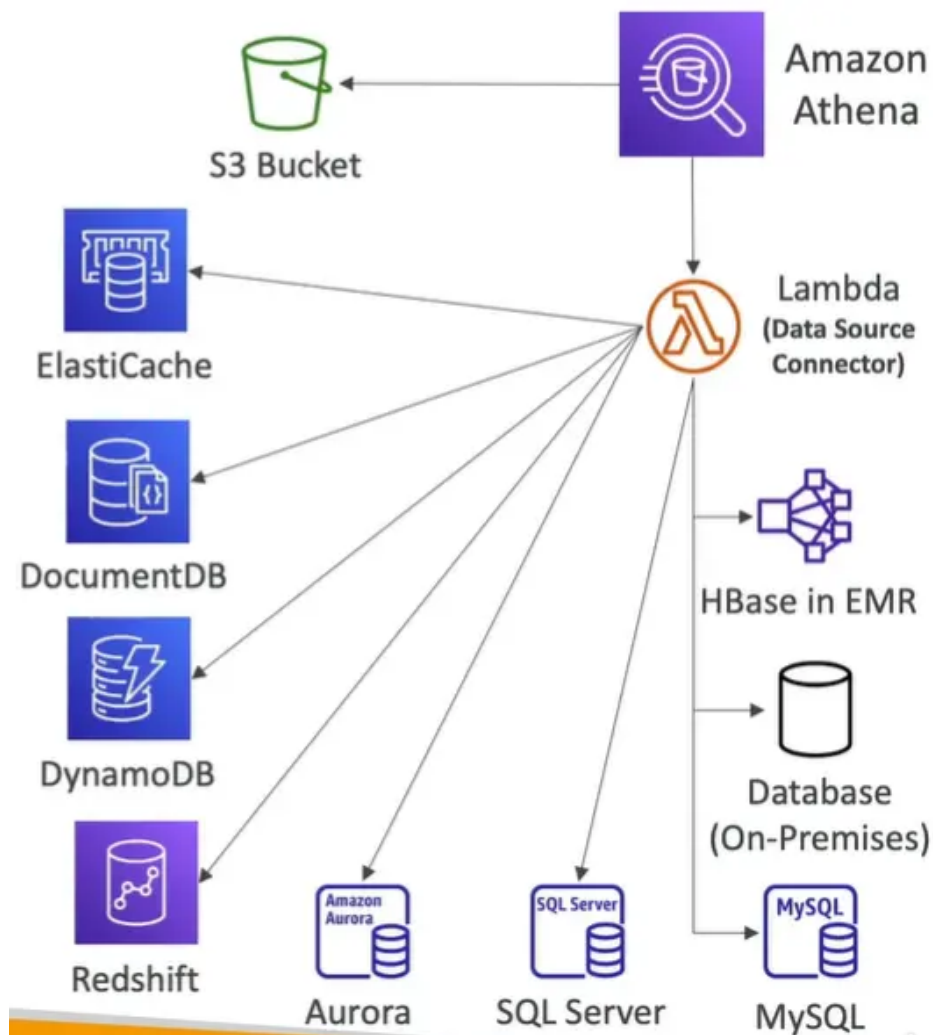
Performance improvement

- use columnar [柱状] data for cost-savings (less scan)
 - apache parquet or ORC is recommended
 - huge performance improvement
 - use Glue to convert your data to Parquet or ORC
- Compress data [压缩数据] for smaller retrievals (bzip2, gzip, lz4, snappy,zip, zstd....)
- Partition datasets in S3 for easy querying on virtual columns


```
s3://yourBucket/pathToTable
    /<PARTITION_COLUMN_NAME>=<VALUE>
    /<PARTITION_COLUMN_NAME>=<VALUE>
    /<PARTITION_COLUMN_NAME>=<VALUE>
    /etc...
```

 - Example: s3://athena-examples/flight/parquet/year=1991/month=1/day=1/
- Use larger files (>128MB) to minimize overhead [开销]

Federated Query [联合查询]



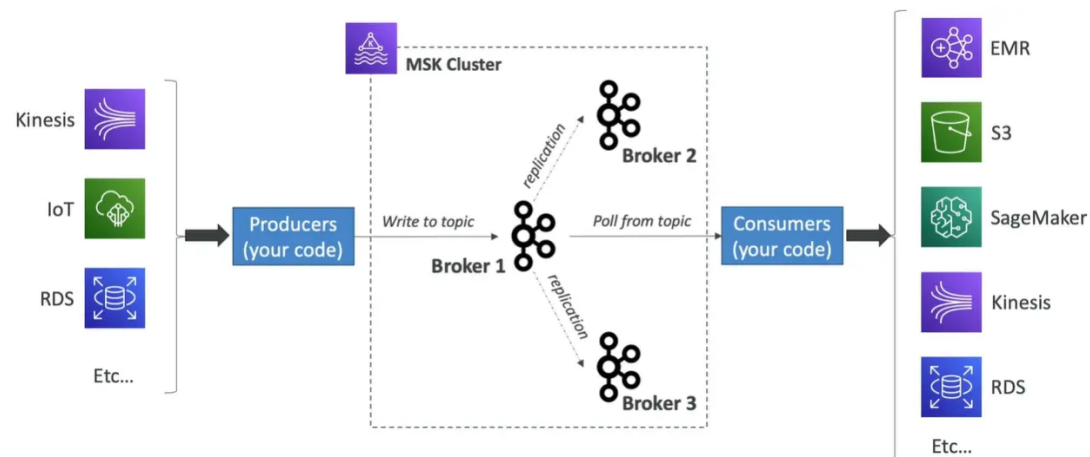
- allows you to run SQL queries across data stored in relational, non-relational, object, and custom data sources (AWS or on-premises)

- use Data Source Connectors that run on AWS Lambda to run Federated Queries (eg, CloudWatch Logs, DynamoSDB, RDS,...)
- store the results back in Amazon S3


Amazon Managed Streaming for Apache Kafka (Amazon MSK)

- Alternative to Amazon Kinesis
- Fully managed Apache Kafka on AWS
 - allow you to create ,update, delete clusters
 - MSK creates & manages Kafka brokers nodes & Zookeeper nodes for you
 - Deploy the MSK cluster in your VPC, multi-AZ (up to 3 for HA)
 - Automatic recovery from common Apache Kafka failures
 - data is stored on EBS volumes for as long as you want
- MSK Serverless
 - run Apache Kafka on MSK without managing the capacity
 - MSK automatically provisions resources and scales compute & storage


Apache Kafka at a high level



Kinesis Data Stream vs, Amazon MSK



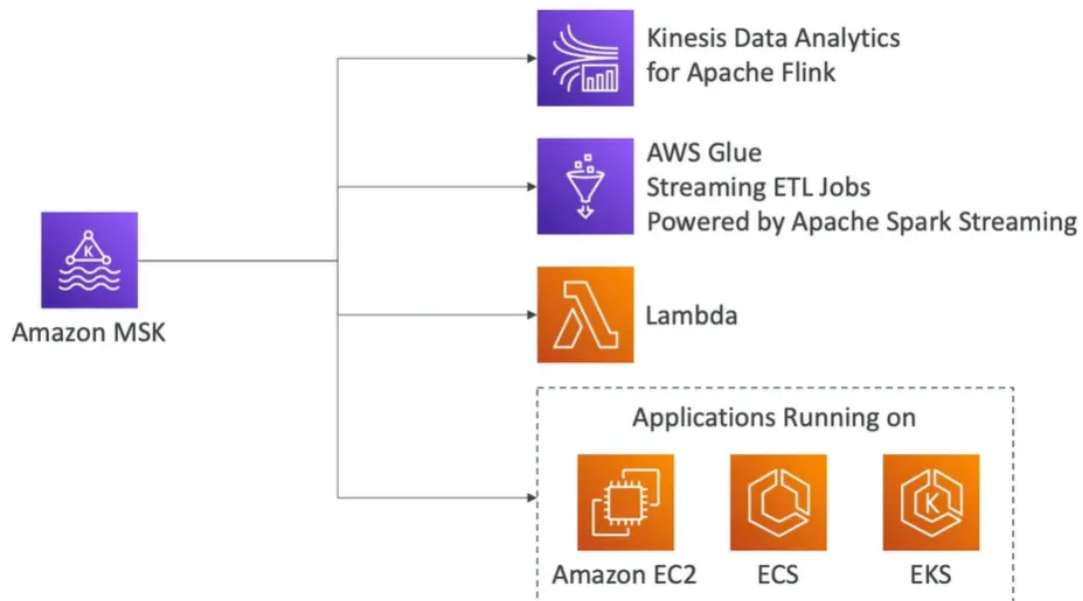
- 1MB message size limit
- data streams with shards
- shard splitting & merging
- TLS in-flight encryption
- KMS at-rest encryption



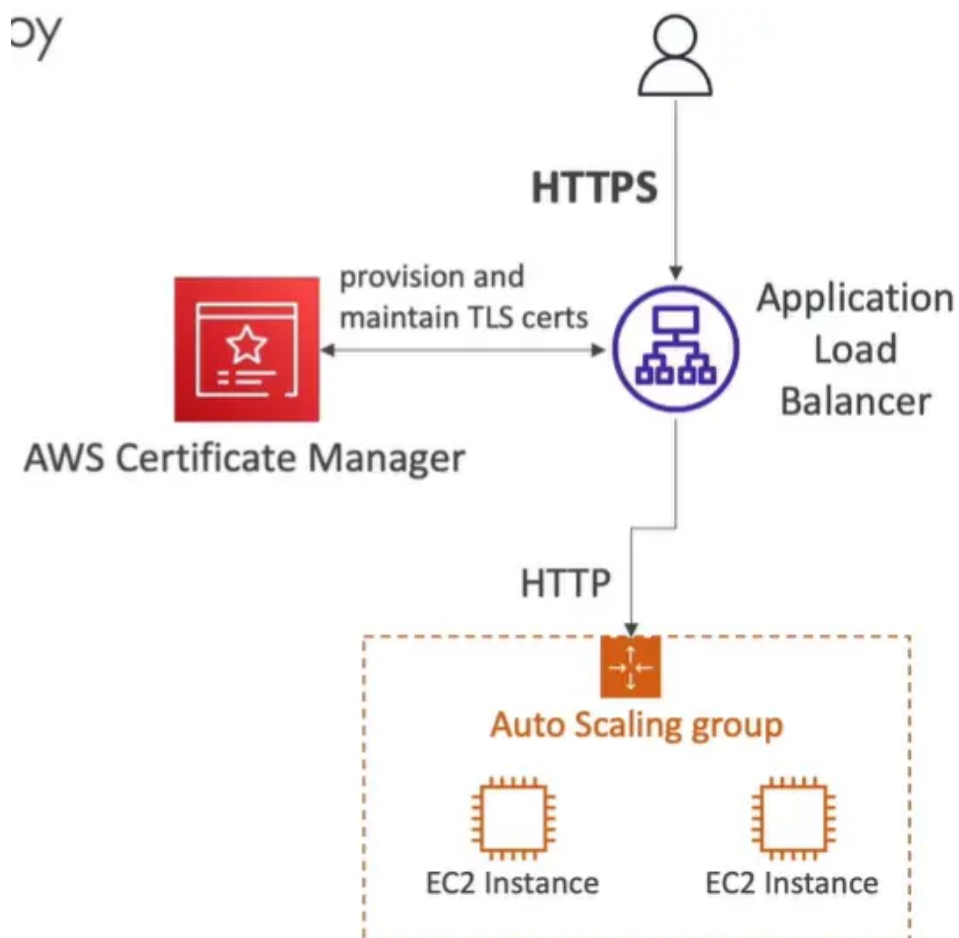
- 1MB default, configure for higher (ex,10MB)
- Kafka Topics with partitions

- can only add partitions to a topic
- PLAINTEXT or TLS In-flight encryption
- KMS at-rest encryption

Amazon MSK Consumers



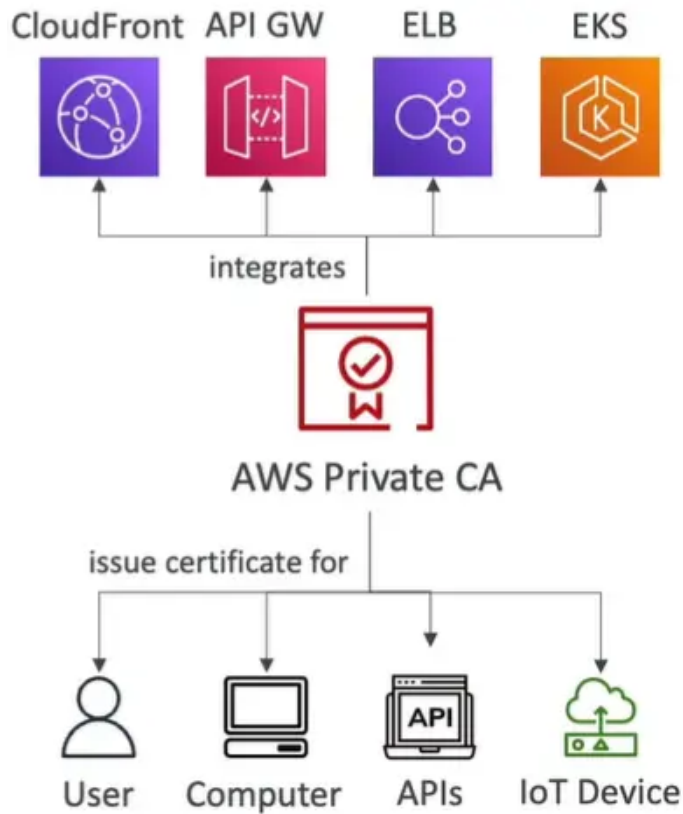
AWS Certificate Manager (ACM)



- let's you easily provision, manage, and deploy SSL / TLS Certificates
- used to provide in-flight encryption for websites (HTTPS)
- supports both public and private TLS certificates
- free of charge for public TLS certificates

- Automatic TLS certificate renewal
- integrations with (load TLS certificate on)
 - Elastic Load Balancers
 - CloudFront Distributions
 - APIs on API Gateway


AWS Private Certificate Authority (CA)



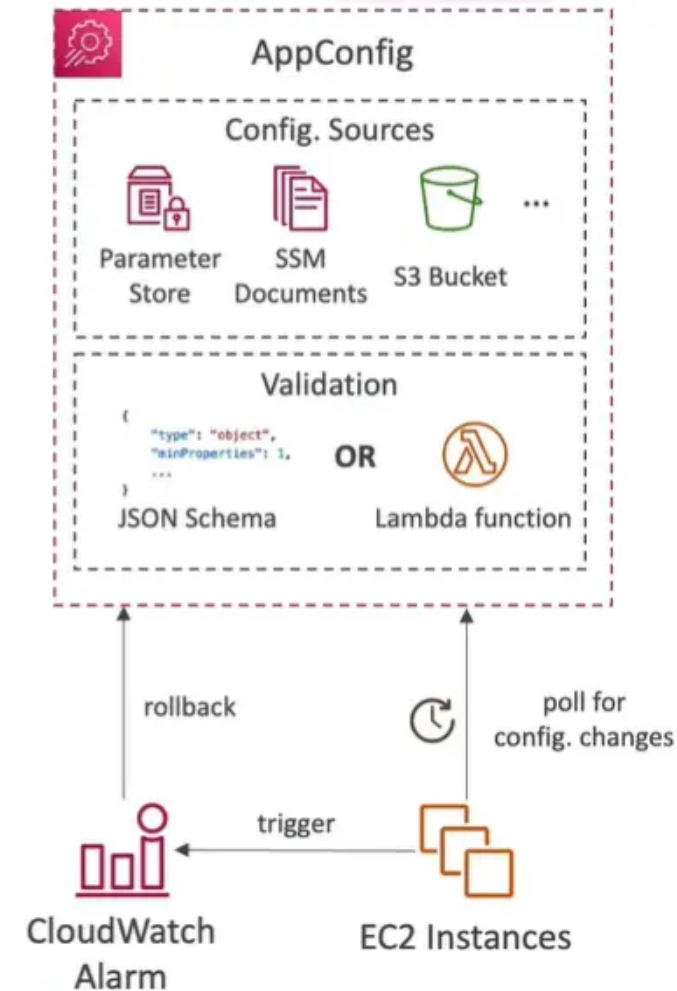
- managed service allows you to create private certificate Authorities (CA), including root and subordinates CAs
- can issue and deploy end-entity X.509 certificates
- certificates are trusted only by your Organization (not the public internet)
- works for AWS service that are integrated with ACM
- use cases:
 - encrypted TLS communication, Cryptographically signing code
 - authenticate users, computers, API endpoints, and IoT devices
 - enterprise customers building a public key infrastructure (PKI)

Amazon Macie

- Amazon Macie is a fully managed data security and data privacy service that uses machine learning and pattern matching to discover and protect your sensitive data in AWS
- Macie help identify and alert you to sensitive data, such as personally identifiable information (PII)

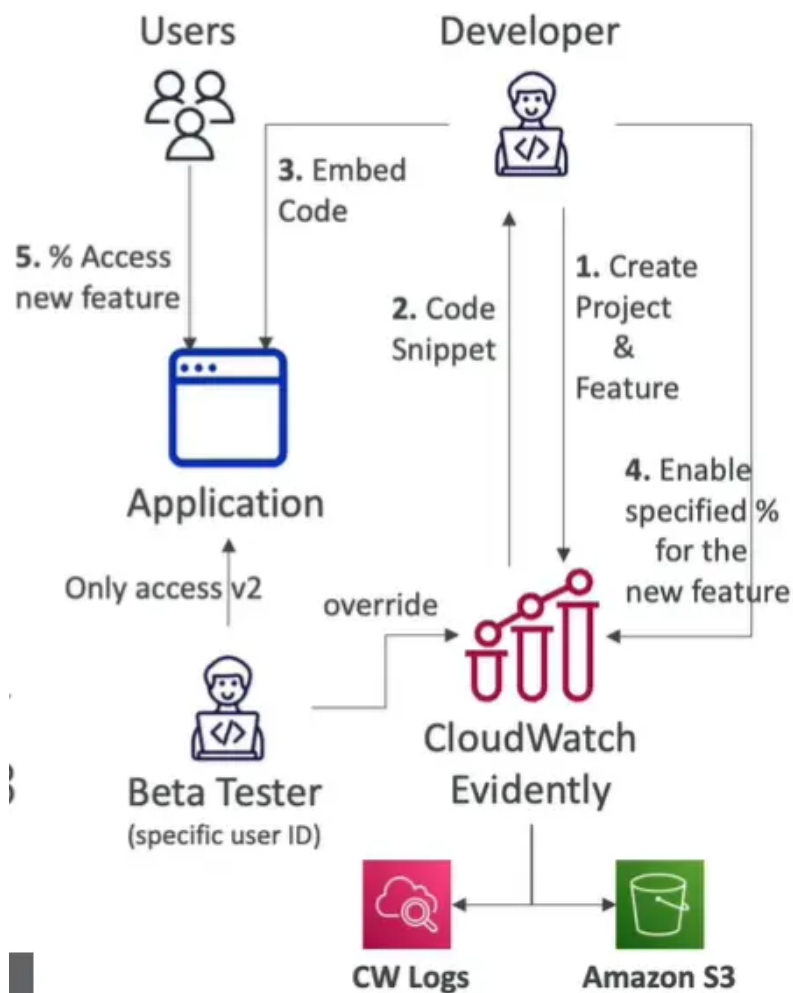


AWS AppConfig



- configure, validate, and deploy dynamic configurations
- deploy dynamic configuration changes to your applications independently of any code deployments
 - you don't need to restart the application
- feature flags ,application tuning, allow/block listing..
- use with apps on EC2 instances, Lambda, ECS, EKS..
- gradually deploy the configuration changes and rollback if issues occur
- validate configuration changes before deployment using
 - JSON Schema (syntactic check) or
 - Lambda Function – run code to perform validation (semantic check)

CloudWatch Evidently



- safely validate new features by serving them to a specified % of your users
 - reduce risk and identify unintended consequences
 - collect experiment data, analyze using stats, monitor performance
- **Launches** (= feature flags) : enable and disable features for a subset of users
- **experiments** (= A/B testing) : compare multiple versions of the same feature
- **overrides** : pre-define a variation for a specific user
- store evaluation events in CloudWatch Logs or S3